AI Report

We are highly confident this text is

AI Generated

AI Probability

100%

This number is the probability that the document is AI generated, not a percentage of AI text in the document.

Plagiarism

?

The plagiarism scan was not run for this document. Go to gptzero.me to check for plagiarism.

A friend at one of the labs told me a couple months back, half kidding, that the only honest pro for - 5/10/2026

A friend at one of the labs told me a couple months back, half kidding, that the only honest pro forma you can write for a frontier model assumes the company dies in three months.
The math: pick Claude N, run its unit economics, watch Claude N+1 kill it, end up sitting on a giant training run that didn't pay back.
Fair.
If that's actually the question, does Claude N pay back as a standalone product, the answer is no, you wouldn't finance that company.
But nobody at the labs writes the pro forma that way, and nobody at Apple writes one for the iPhone 12 either.
The unit isn't the SKU.
It's the line.
The current model earns revenue, the previous one becomes the cheap routing tier, this generation's failed trajectories become next generation's training data, the integrations carry forward, the serving stack gets cheaper (Epoch has inference cost halving every two months at fixed capability, a real input the 2024 pieces did not price in), the contracts compound, and the only thing that ages out in three months is the checkpoint itself.
Traces don t. Contracts don t. FedRAMP doesn t. The question is whether the rolling machine throws off cash while it keeps moving.

OK granting that which resets the unit of analysis more than answers anything where does the margin actually come from?
Best guess the moment a lab can charge for the work an agent does instead of the tokens it spends doing it.
Chat was a commodity from week one.
It competed against search against interns against mild conversational convenience all of which have held the consumer subscription ceiling at 20 month since 2023.
Agents are different.
A coding agent at 4 per run that saves three engineer hours does not get priced against the 4.

It gets priced against a fully loaded 250 an hour engineer and at that point arguing about token margin gets silly.
A finance agent that catches a bad assumption inside a deal model clears against the deal it kept.
SemiAnalysis writes the inequality out cleanly in the value capture piece customers pay more for successful trajectories labs spend each quarter making those trajectories cheaper to produce caching kernels hardware and post training compression nobody likes calling compression and the lab pockets the spread.
The spread is huge when quality has cliffs.
A model 5 behind on a benchmark produces 50 fewer successful long agentic runs because errors compound.
One hallucinated API call.
No partial credit.

The obvious counter is that open source eats this the way it ate everything else and until pretty recently I thought it would.
Recipes were shared everyone read the same papers the same dozen researchers cycled labs every eighteen months the public benchmarks defined the frontier.
DeepSeek s whole pitch was take the tricks add taste and ruthless engineering reach the frontier on a fraction of the budget.
That worked.
What changed in the last year is the recipes have gone proprietary.
Aschenbrenner has been writing about this from inside the sober version not the millenarian part algorithmic ideas diverging the data wall biting some labs and not others agentic curricula now lab specific.
NIST s CAISI evaluation of DeepSeek V4 Pro this month is the cleanest public confirmation I ve seen.
DeepSeek looks competitive on the public numbers everyone quotes.
On CAISI s private agentic and cyber evals it lands about eight months back the kind of gap that matters when customers are buying finished work.
In that regime open source takes the boring tier summarization classification autocomplete batch transforms and the labs take what breaks without the frontier model.
Different businesses.

Compute is the piece the early labs can t make money pieces ignored.
If any single lab had the capacity to serve the entire agentic market it would have to underprice itself against the next best alternative the standard story for API pricing converging.
But every frontier lab is capacity constrained on chips on power on the data center buildout itself and that s going to be true through 2030 unless someone has a credible chip glut forecast I have not seen.
When three or four labs sell out at once which is roughly the state of the world right now prices clear against customer ROI rather than against the marginal open source token.
The labs only print money after the scaling plateau framing has this exactly inverted.
The plateau is when supply catches demand and margins compress.
The race is the high margin period.
Labs make their money during the race they keep some after.

The deepest worry in the original question is about switching costs.
Fair.
Cloud margins come from switching costs and an API endpoint is easy to swap.
What that misses is that the product isn t really the API anymore.
It s the agent.

An agent wired into your codebase your CRM your data warehouse your Snowflake instance your ticket queue your audit logs and your legal archive is not portable in any meaningful sense.
The customer spent four months on security review evals permissions change management training and internal trust.
Replacing the agent puts a working process at risk and large customers will pay quite a lot of money to not do that.
The failure mode I do worry about and I want to be honest I haven t talked myself out of it is the one where open models close to within a month of the frontier on private agentic benchmarks at the same moment a chip glut arrives.
Call that 20 not zero.
In the other 80 the gap stays at six to twelve months on the benchmarks customers pay against and the labs don t need any single model to pay itself back.
They need to be the bottleneck for work customers can t afford to botch.
Easier bar than earning a training run back in a quarter and from what I can tell they re already over it.

High Human Impact High AI Impact

FAQs

What is GPTZero?

GPTZero is the leading AI detector for checking whether a document was written by a large language model such as ChatGPT. GPTZero detects AI on sentence, paragraph, and document level. Our model was trained on a large, diverse corpus of human-written and AI-generated text with support for English, Spanish, French, German, and other languages. To date, GPTZero has served over 17 million users around the world, and works with over 100 organizations in education, hiring, publishing, legal, and more.

When should I use GPTZero?

Our users have seen the use of AI-generated text proliferate into education, certification, hiring and recruitment, social writing platforms, disinformation, and beyond. We've created GPTZero as a tool to highlight the possible use of AI in writing text. In particular, we focus on classifying AI use in prose. Overall, our classifier is intended to be used to flag situations in which a conversation can be started (for example, between educators and students) to drive further inquiry and spread awareness of the risks of using AI in written work.

Does GPTZero only detect ChatGPT outputs?

No, GPTZero works robustly across a range of AI language models, including but not limited to ChatGPT, GPT-5, GPT-4, GPT-3, Gemini, Claude, and AI services based on those models.

What are the limitations of the classifier?

The nature of AI-generated content is changing constantly. As such, these results should not be used to punish students. We recommend educators to use our behind-the-scene Writing Reports as part of a holistic assessment of student work. There always exist edge cases with both instances where AI is classified as human, and human is classified as AI. Instead, we recommend educators take approaches that give students the opportunity to demonstrate their understanding in a controlled environment and craft assignments that cannot be solved with AI. Our classifier is not trained to identify AI-generated text after it has been heavily modified after generation (although we estimate this is a minority of the uses for AI-generation at the moment). Currently, our classifier can sometimes flag other machine-generated or highly procedural text as AI-generated, and as such, should be used on more descriptive portions of text.

I'm an educator who has found AI-generated text by my students. What do I do?

Firstly, at GPTZero, we don't believe that any AI detector is perfect. There always exist edge cases with both instances where AI is classified as human, and human is classified as AI. Nonetheless, we recommend that educators can do the following when they get a positive detection: Ask students to demonstrate their understanding in a controlled environment, whether that is through an in-person assessment, or through an editor that can track their edit history (for instance, using our Writing Reports through Google Docs). Check out our list of several recommendations on types of assignments that are difficult to solve with AI.

Ask the student if they can produce artifacts of their writing process, whether it is drafts, revision histories, or brainstorming notes. For example, if the editor they used to write the text has an edit history (such as Google Docs), and it was typed out with several edits over a reasonable period of time, it is likely the student work is authentic. You can use GPTZero's Writing Reports to replay the student's writing process, and view signals that indicate the authenticity of the work.

See if there is a history of AI-generated text in the student's work. We recommend looking for a long-term pattern of AI use, as opposed to a single instance, in order to determine whether the student is using AI.